

Machine learning and physical modelling-2

julien.brajard@nersc.no

October 2019

NERSC

<https://github.com/brajard/MAT330>

1. Steps of a machine learning process
2. A standard Machine learning model: Random Forests
3. Neural Networks
4. A quick typology of few neural nets


Steps of a machine learning process

Steps

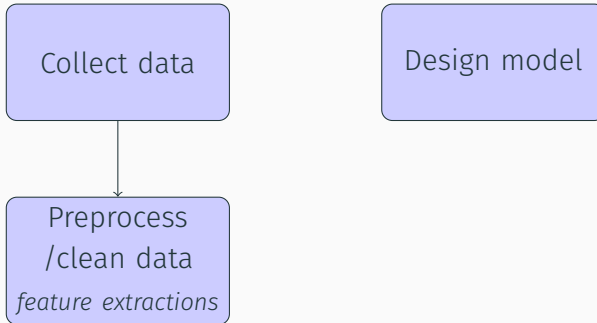


Collect data


Steps



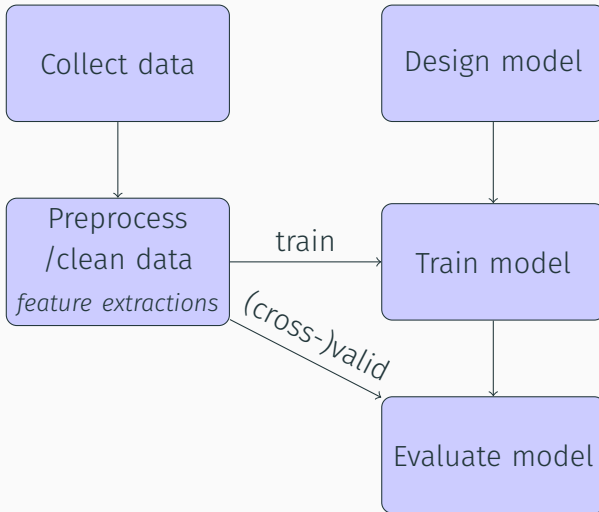
Steps



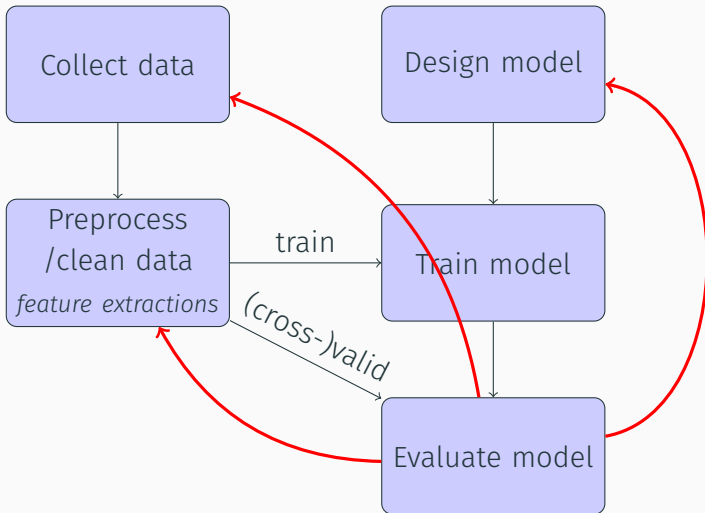
Steps



Steps



Steps



In summary

From one dataset, 3 sub-datasets have to be extracted:

- A training dataset
- A validation dataset

Can be done iteratively in a cross-validation procedure.

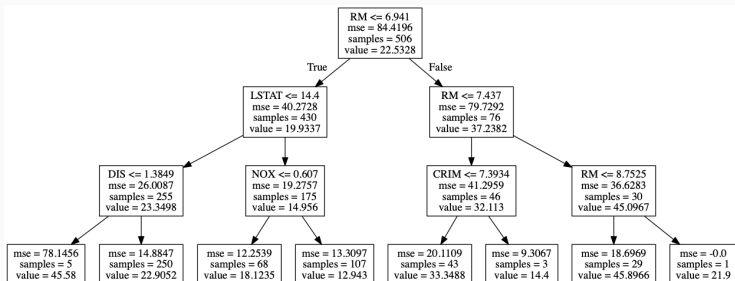
Some parameters of the model (e.g. polynomial order in a polynomial regression) were determined from the validation dataset.

- A test dataset (independent from the two other) to estimate the final performance of the model.

A standard Machine learning model:
Random Forests


A decision tree

Predict house price (in \$1000's) from 13 features:





CRIM	per capita crime rate by town
NOX	nitric oxides concentration
RM	average number of rooms per dwelling
DIS	distance to employment centres
LSTAT	lower status of the population



Uni-variate example



Uni-variate example



Uni-variate example



From tree to forest

Disadvantages of regression tree:

- Can overfit the data



From tree to forest

Disadvantages of regression tree:


- Can overfit the data

One extension of Regression Tree: Random Forest




The (over simplified) principle of Random Forest


Data




The (over simplified) principle of Random Forest




The (over simplified) principle of Random Forest




The (over simplified) principle of Random Forest



The (over simplified) principle of Random Forest

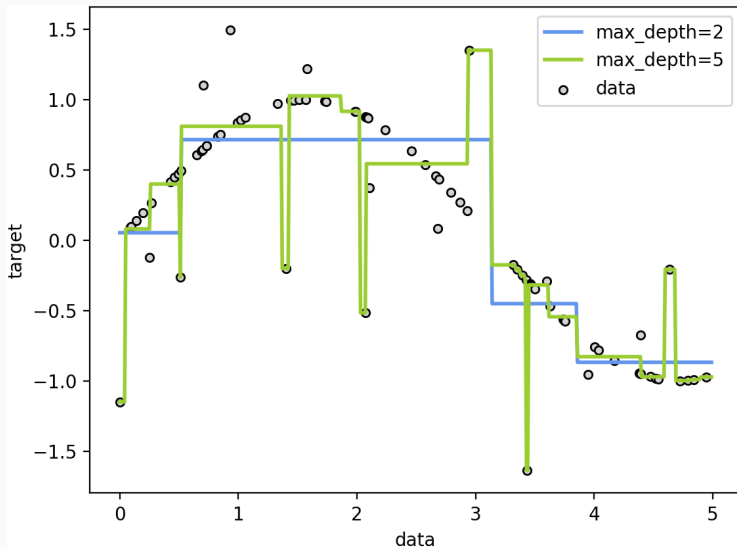


The (over simplified) principle of Random Forest




Results on the univariate experiment

Prediction of Randoms trees



Results on the univariate experiment

Prediction of a Random Forest



Some key parameters

```
from sklearn.ensemble import RandomForestClassifier  
  
rf = RandomForestRegressor(n_estimators=n, max_features=  
    maxf, min_samples_split=min_split, ...)
```

- **n_estimators**: number of trees (generally the larger is the better)

Some key parameters

```
from sklearn.ensemble import RandomForestClassifier  
  
rf = RandomForestRegressor(n_estimators=n, max_features=  
    maxf, min_samples_split=min_split, ...)
```

- **n_estimators**: number of trees (generally the larger is the better)
- **max_features**: number of features to consider at each split. The default number is the total number of features. A larger value makes provides a smaller bias (accuracy) but a bigger variance (risk of overfitting)

Some key parameters


```
from sklearn.ensemble import RandomForestClassifier  
  
rf = RandomForestRegressor(n_estimators=n, max_features=  
    maxf, min_samples_split=min_split, ...)
```

- **n_estimators**: number of trees (generally the larger is the better)
- **max_features**: number of features to consider at each split. The default number is the total number of features. A larger value makes provides a smaller bias (accuracy) but a bigger variance (risk of overfitting)
- **min_samples_fit**: number of features to consider at each split. The minimum value of 2 means that the tree is fully developed (small bias but great variance).

Feature importance

```
rf = RandomForestRegressor(n_estimators=1000,  
    max_features=10,random_state=10)  
rf.fit(X,y)  
importances = rf.feature_importances_
```

Indicates the impact of a feature in predicting the target.



CRIM	per capita crime rate by town
NOX	nitric oxides concentration
RM	average number of rooms per dwelling
DIS	distance to employment centres
LSTAT	lower status of the population

Determination of the hyperparameters

- Parameters that are not optimized during the training are called **hyperparameters**.

Determination of the hyperparameters


- Parameters that are not optimized during the training are called **hyperparameters**.
- They can be determined using a score on the **validation dataset** or using a **cross-validation** procedure.

Determination of the hyperparameters

- Parameters that are not optimized during the training are called **hyperparameters**.
- They can be determined using a score on the **validation dataset** or using a **cross-validation** procedure.

Determination of the hyperparameters


- Parameters that are not optimized during the training are called **hyperparameters**.
- They can be determined using a score on the **validation dataset** or using a **cross-validation** procedure.



<https://xkcd.com/1838/>

Stir the pile: The gridsearch

1. Specify a list of hyperparameters to be tested.
2. For each of the parameters, specify a set of values to test
3. Train a model for each of the possible combinations of hyperparameters
4. Retain the best model (using, e.g., cross-validation)



[https://medium.com/
@senapati.dipak97/grid-
search-vs-random-search-
d34c92946318](https://medium.com/@senapati.dipak97/grid-search-vs-random-search-d34c92946318)

Remarks on the gridsearch procedure

- It make an **exhaustive** search of the hyperparameters

Remarks on the gridsearch procedure

- It make an **exhaustive** search of the hyperparameters
- The procedure is easy to **parallelized**.

Remarks on the gridsearch procedure


- It make an **exhaustive** search of the hyperparameters
- The procedure is easy to **parallelized**.
- it is not naturally adapted for quantitative hyperparameters.

Remarks on the gridsearch procedure

- It make an **exhaustive** search of the hyperparameters
- The procedure is easy to **parallelized**.
- it is not naturally adapted for quantitative hyperparameters.
- it can become **very costly**. (e.g. 8 hyperparameters with 8 values each to test = $8^8 = 16,777,216$ trainings).

Random search

1. Specify a list of hyperparameters to be tested.
2. For each of the parameters, specify a set of values to test or a law to draw a random value.
3. Draw n combinations of the hyperparameters.
4. Train a model for each of the combinations.
5. Retain the best model (using, e.g., cross-validation)



Remarks on the random search procedure

- It **does not make** an **exhaustive** search of the hyperparameters
- The procedure is easy to **parallelized**.
- it is not adapted for quantitative hyperparameters.
- The cost is predictable (number of draw).


Remarks on the random search procedure

- It **does not make** an **exhaustive** search of the hyperparameters
- The procedure is easy to **parallelized**.
- it is not adapted for quantitative hyperparameters.
- The cost is predictable (number of draw).

Both gridsearch and random search are implemented and easy to use in scikit-learn.

Neural Networks

The perceptron : an artificial neuron



Computation

$$y = f(w_0 + w_1.x_1 + w_2.x_2 + \cdots + w_n.x_n) = f(w_0 + \sum_{i=1}^n w_i.x_i)$$

Some remarks

- Inputs x_i are the different features of the data

Some remarks

- Inputs x_i are the different features of the data
- Weight w_i are the parameters of the model to optimize

Some remarks


- Inputs x_i are the different features of the data
- Weight w_i are the parameters of the model to optimize
- If the activation function is identity, it is equivalent to a linear regression

Some remarks

- Inputs x_i are the different features of the data
- Weight w_i are the parameters of the model to optimize
- If the activation function is identity, it is equivalent to a linear regression


More complexe models are build by combining several perceptrons

Multi-layer perceptron (Densely connected layers)




Most usual activation functions


Linear




Hyperbolic tangent



ReLU




Sigmoid




New fancy activation functions

Leaky-ReLU



Swish



Regression

- Last layer:
linear or hyperbolic
tangent
- Loss function:

$$L(\hat{y}, y) = \sum_i (\hat{y}_i - y_i)^2$$

Classification and regression loss

Regression

- Last layer:
linear or hyperbolic
tangent
- Loss function:

$$L(\hat{y}, y) = \sum_i (\hat{y}_i - y_i)^2$$

Classification

- Last layer:
Soft-max

$$p_j = f_j(\mathbf{h}) = \frac{e^{h_j}}{\sum_k e^{h_k}}$$

- Loss function:
Negative crossentropy

$$L(p, y) = - \sum_i \sum_j y_{i,j} \cdot \log p_{i,j}$$

Convolutional neural net

X: an image

x_{11}	x_{12}	x_{13}	x_{14}	x_{15}	x_{16}
x_{21}	x_{22}	x_{23}	x_{24}	x_{25}	x_{26}
x_{31}	x_{32}	x_{33}	x_{34}	x_{35}	x_{36}
x_{41}	x_{42}	x_{43}	x_{44}	x_{45}	x_{46}
x_{51}	x_{52}	x_{53}	x_{54}	x_{55}	x_{56}
x_{61}	x_{62}	x_{63}	x_{64}	x_{65}	x_{66}

$w_{11} w_{12} w_{13}$

$w_{21} w_{22} w_{23}$

$w_{31} w_{32} w_{33}$

W



h: first feature

h_{11}	h_{12}	h_{13}	h_{14}
h_{21}	h_{22}	h_{23}	h_{24}
h_{31}	h_{32}	h_{33}	h_{34}
h_{41}	h_{42}	h_{43}	h_{44}

Perform a standard convolution

$$h_{i,j} = \sum_{k=1}^3 \sum_{l=1}^3 x_{i+k-1,j+l-1} \cdot w_{k,l}$$


Main parameters of a convolutional layer

- Size of the filter K

Main parameters of a convolutional layer

- Size of the filter K
- Number of filters p

A convolutional layer is composed of p convolutions (size of layer) extracting p features from the data.




$$O = \frac{W-K+2P}{S} + 1, \text{ where } O \text{ is the output size and } W \text{ the input size.}$$

Main parameters of a convolutional layer


- Size of the filter K
- Number of filters p
- Strides S

$S = 1$

7 x 7 Input Volume




5 x 5 Output Volume



$S = 2$

7 x 7 Input Volume



3 x 3 Output Volume

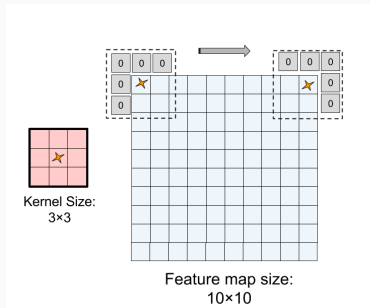


$$O = \frac{W-K+2P}{S} + 1, \text{ where } O \text{ is the output size and } W \text{ the input size.}$$

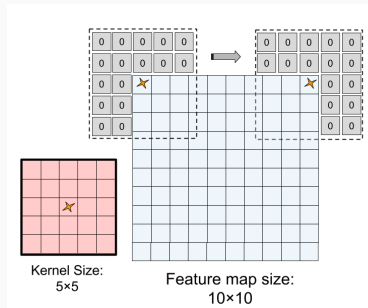
Main parameters of a convolutional layer

- Size of the filter K
- Number of filters p
- Strides S
- Padding P

$P = 1$






$P = 2$



$$O = \frac{W - K + 2P}{S} + 1, \text{ where } O \text{ is the output size and } W \text{ the input size.}$$




Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution






Summary of Convolutional layer steps

1. Convolution




Summary of Convolutional layer steps

1. Convolution



Summary of Convolutional layer steps

1. Convolution




2. Addition




Summary of Convolutional layer steps

1. Convolution




2. Addition




Summary of Convolutional layer steps

1. Convolution




2. Addition




Summary of Convolutional layer steps

1. Convolution




2. Addition




Summary of Convolutional layer steps

1. Convolution



2. Addition




3. Bias




Summary of Convolutional layer steps


1. Convolution



2. Addition




3. Bias




Summary of Convolutional layer steps


1. Convolution



2. Addition



3. Bias




Remarks on Convolutional layers

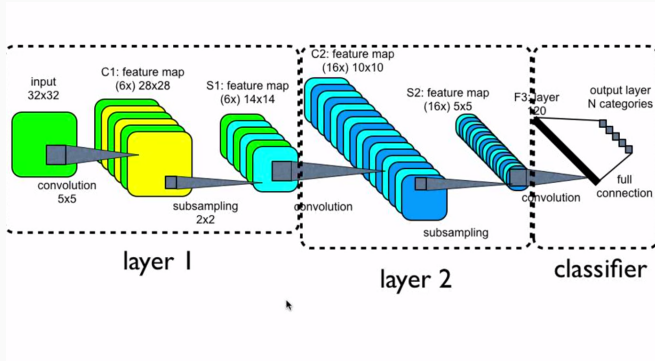
- Convolutional layers are acting locally on the image (But you can still use large scale information by adding more layers)
- Convolutions are invariant by translation (the weights do not depend on the location on the image).
- They can handle images of different sizes.

Max-Pooling

In order to reduce the size of the feature space (en to enhance the gradients), a common operation is to perform a max-pooling.



A traditional CNN architecture



Example of AlexNet

AlexNet is the first Deep architecture used on ImageNet challenge in 2012 and achieved an **error of 15.3%** (10% better than the previous best classifier). The paper was cited more than 34,000 times.




Alex Krizhevsky and Geoffrey E Hinton, *ImageNet Classification with Deep Convolutional Neural Networks*, Neural Information Processing Systems (2012), 1–9.

Layer		Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	227x227x3	-	-	-
1	Convolution	96	55 x 55 x 96	11x11	4	relu
	Max Pooling	96	27 x 27 x 96	3x3	2	relu
2	Convolution	256	27 x 27 x 256	5x5	1	relu
	Max Pooling	256	13 x 13 x 256	3x3	2	relu
3	Convolution	384	13 x 13 x 384	3x3	1	relu
4	Convolution	384	13 x 13 x 384	3x3	1	relu
5	Convolution	256	13 x 13 x 256	3x3	1	relu
	Max Pooling	256	6 x 6 x 256	3x3	2	relu
6	FC	-	9216	-	-	relu
7	FC	-	4096	-	-	relu
8	FC	-	4096	-	-	relu
Output	FC	-	1000	-	-	Softmax

A quick typology of few neural nets

Recurrent Neural Networks




Some popular types of recurrent neural networks:

- Long short-term memory (LSTM)
- Gated Recurrent Unit (GRU)

Used in machine translation and text processing


Autoencoders



Used in image denoising, compressing, generation,...


Generative adversarial networks

Generative adversarial networks (conceptual)



5


Residual Networks



x : input, y : output


$$y = x + \mathcal{F}(x)$$

Residual Networks



x : input, y : output

$$y = x + \mathcal{F}(x)$$



Questions addressed in this lecture

- What are the steps of a machine learning process?
- What is the principle of the Random Forests? [Van16,5.8]
- How to determine the hyperparameters? [Van16,5.3]
- What is an artificial neural network? [GBC16,6]
- What is a convolutional layer? [GBC16,9]
- What are the main types of neural networks? [GBC16,10, 14, 20]

<https://towardsdatascience.com/intuitively-understanding-convolutions-for-deep-learning-1f6f42faee1>

Refs

[Van16,*n*]: Jake VanderPlas, *Python Data Science Handbook*, section *n*
[GBC16,*n*]: Goodfellow et al., *Deep Learning*, chapter *n*